

Lecture AI3611:智能感知认知实践

05 实践项目

Mengyue Wu & Xie Chen

Cross Media Language Intelligence Lab (X-LANCE)
Department of Computer Science & Engineering
Shanghai Jiao Tong University

2024

- ▶ 单模态理解
 - ▶ 基于Vall-E模型的语音合成
 - ▶ 声音事件检测
 - ▶ 语言模型
 - ▶ 图像生成
- ▶ 多模态及跨模态交互
 - ▶ 图片摘要生成
 - ▶ 音视频场景识别

6选4进入计分, 可以自由选择

- ▶ 单模态理解
 - ▶ 基于Vall-E模型的语音合成
 - ▶ 声音事件检测
 - ▶ 语言模型
 - ▶ 图像生成
- ▶ 多模态及跨模态交互
 - ▶ 图片摘要生成
 - ▶ 音视频场景识别
- ▶ 实践项目单个21%, 6选4, 共84%
 - ▶ 按兴趣选择
 - ▶ 按自我最优提交

- ▶ 实践考核

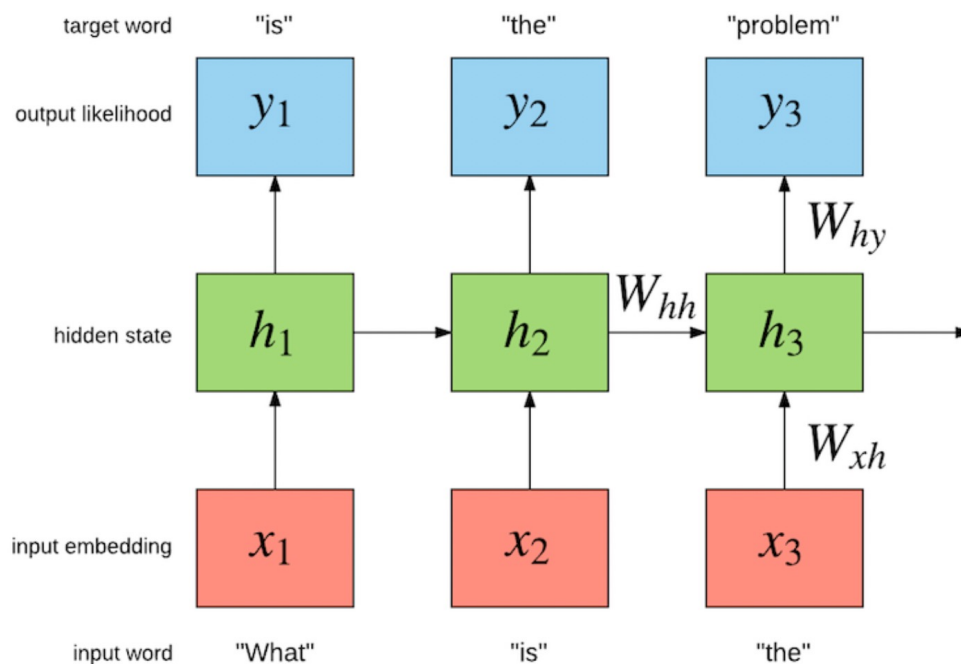
- ▶ 1. 项目报告 - 40% 清楚阐述实验过程，实验结果分析
- ▶ 2. 模型性能 - 40% 模型创新程度，测试集性能
- ▶ 3. 代码可读性 - 20% 逻辑清晰，易复现，注释等

- ▶ **语言模型**
- ▶ 语音合成
- ▶ 图像生成
- ▶ 声音事件检测
- ▶ 图片摘要生成
- ▶ 音视频场景识别

► Perplexity (PPL)

- the lower the better
- average divergence of each prediction

$$PPL = \exp\left(-\frac{1}{N} \sum_{i=1}^N \log P(w_i|h)\right)$$



► 基本要求

- 使用不同模型结构，训练基于神经网络的语言模型（至少尝试Feedforward, RNN, GRU, LSTM, Transformer等结构中的三种不同神经网络模型结构）
- 讨论和尝试不同超参数，对语言模型性能的影响
- 模型总参数量不得超过60M，最优化测试集上的PPL

► 实践报告要求

1. 回答基本要求中的各个要点
2. 写出详细的实验过程和实验分析
3. 提交代码，给出重现最优结果的脚本和配置

► 加分项目

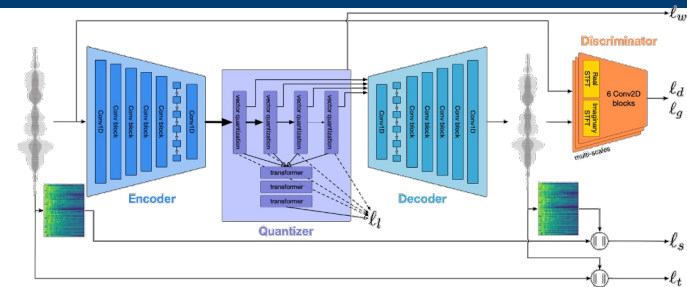
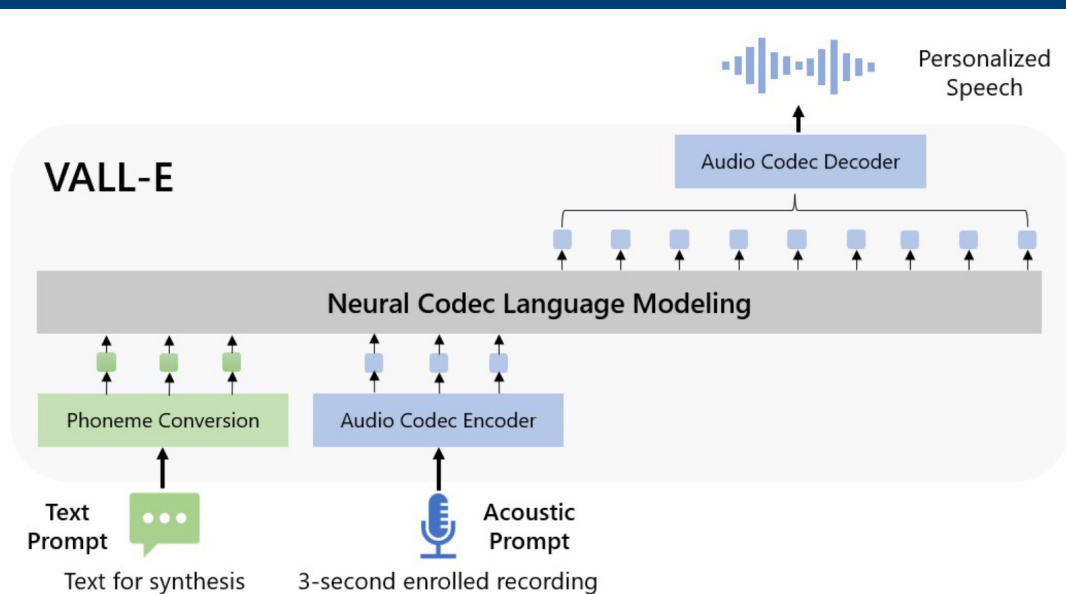
- 使用tensorboard画出训练阶段每个epoch的train, valid和test loss(或PPL) 趋势
- 解释和解决Transformer LM性能不如LSTM LM的问题
- 可以使用相关文献中或Github中开放的更先进的语言模型，对本课程提供的数据持续提升性能

▶ 参考文献

- ▶ Mikolov, Tomas, Martin Karafiát, Lukas Burget, Jan Cernocký, and Sanjeev Khudanpur. "Recurrent neural network based language model." In *Interspeech*, vol. 2, no. 3, pp. 1045-1048. 2010.
- ▶ Sundermeyer, Martin, Ralf Schlüter, and Hermann Ney. "LSTM neural networks for language modeling." In *Thirteenth annual conference of the international speech communication association*. 2012.
- ▶ Irie, Kazuki, Albert Zeyer, Ralf Schlüter, and Hermann Ney. "Language modeling with deep transformers." *arXiv preprint arXiv:1905.04226* (2019).

- ▶ 语言模型
- ▶ **语音合成**
- ▶ 图像生成
- ▶ 声音事件检测
- ▶ 图片摘要生成
- ▶ 音视频场景识别

语音合成 – Vall-E模型



- **This outward mutability** indicated, and did not more than fairly express, the various properties of her inner life.



Prompt



VALL-E

- **This she said was true hospitality**, and I am not sure that I did not agree with her.



Prompt



VALL-E

► 基本要求

1. 基本理解代码逻辑，熟悉基于Vall-E语音合成过程
2. 得到不同训练数据规模下（100，360和960小时）的语音合成性能对比
3. 各模型参数总数量不超过500M（现有例子模型参数已超过60M）
4. 评价指标：最小化词错误率（WER）；说话人相似性
 1. 基于已有的语音识别系统评测识别性能
 2. 基于已有的说话人识别模型评价说话人相似性

► 实践报告要求

1. 回答基本要求中的各个要点
2. 写出详细的实验过程和实验分析
3. 提交代码，给出重现最优结果的脚本和配置

► 加分项

- 尝试使用已有的中文语音数据，类似思路搭建一个中文语音合成系统
- 尝试不同的neural codec模型，提高合成语音质量
- 尝试不同模型结构或算法，提高模型稳定性或语音质量


注：该题目对算力有一定要求，使用学校集群，可能会导致超过最长时长，需要断点

▶ 参考文献

- ▶ Défossez, Alexandre, Jade Copet, Gabriel Synnaeve, and Yossi Adi. "High fidelity neural audio compression." *arXiv preprint arXiv:2210.13438* (2022).
- ▶ Wang, Chengyi, et al. "Neural codec language models are zero-shot text to speech synthesizers." *arXiv preprint arXiv:2301.02111* (2023).
- ▶ Kumar, Rithesh, Prem Seetharaman, Alejandro Luebs, Ishaan Kumar, and Kundan Kumar. "High-fidelity audio compression with improved rvqgan." *Advances in Neural Information Processing Systems* 36 (2024).

- ▶ 语言模型
- ▶ 语音合成
- ▶ **图像生成**
- ▶ 声音事件检测
- ▶ 图片摘要生成
- ▶ 音视频场景识别

图像生成 – Variational AutoEncoders (VAE)



- ▶ Loss function consists of two parts
 - ▶ regularization loss: KL-distance between $N(\mu, \sigma)$ and $N(0, I)$
 - ▶ reconstruction loss: reconstruct the image in the output layer of decoder
- ▶ Notes and tips:
 - ▶ KL distance between two Gaussian distributions: $KLD(p, q) = \log \frac{\sigma_2}{\sigma_1} + \frac{\sigma_1^2 + (u_1 - u_2)^2}{2\sigma_2^2}$
 - ▶ The error (or gradient) not back propagation from sampling (dotted line)
 - ▶ The **sampled latent vector** used the μ, σ and samples from $N(0, I)$ to generate the input for decoder. only μ, σ connection used for error backpropagation
 - ▶ value in MNIST used here are binary value 0 or 1, consider using BCELoss in Pytorch

▶ 自由发挥题

- ▶ 提供MNIST数据集，和data loader代码，不提供模型训练样本代码
- ▶ 不限模型结构，模型大小
- ▶ 可以借鉴Github代码，但不要照搬，发现N份雷同代码($N > 1$)，每份雷同代码作业扣N分


▶ 要求和目标

- ▶ 给定MNIST数据集 (train和valid)，自定义encoder和decoder，实现VAE训练
- ▶ 将隐层向量z维度设置为1，比较VAE训练完成后不同的z值对应的生成图片效果
- ▶ 将隐层向量z维度设置为2，找出隐层向量的两个维度 $[-5, 5]$ 值区间内对应的图片生成效果
- ▶ 最优化valid数据集的重构误差


▶ 报告要求

- ▶ 根据自己的理解，描述VAE的原理，可用数学公式，或图片，或文字
- ▶ 写出详细的实验过程和实验分析
- ▶ 提交代码（需包含代码注释，方便阅读），给出重现最优结果的脚本和配置
- ▶ 鼓励加上你觉得这个模型有意思的观察或思考！

2-dim gaussian



1-dim gaussian



▶ 参考文献

- ▶ Kingma, Diederik P., and Max Welling. "Auto-encoding variational bayes." *arXiv preprint arXiv:1312.6114* (2013).
- ▶ Doersch, Carl. "Tutorial on variational autoencoders." *arXiv preprint arXiv:1606.05908* (2016).

课程作业问题联系人

授课教师： 陈谐 邮箱： chenxie95@sjtu.edu.cn
助教： 马子阳 邮箱： zym.22@sjtu.edu.cn

微信可从课程微信群里添加

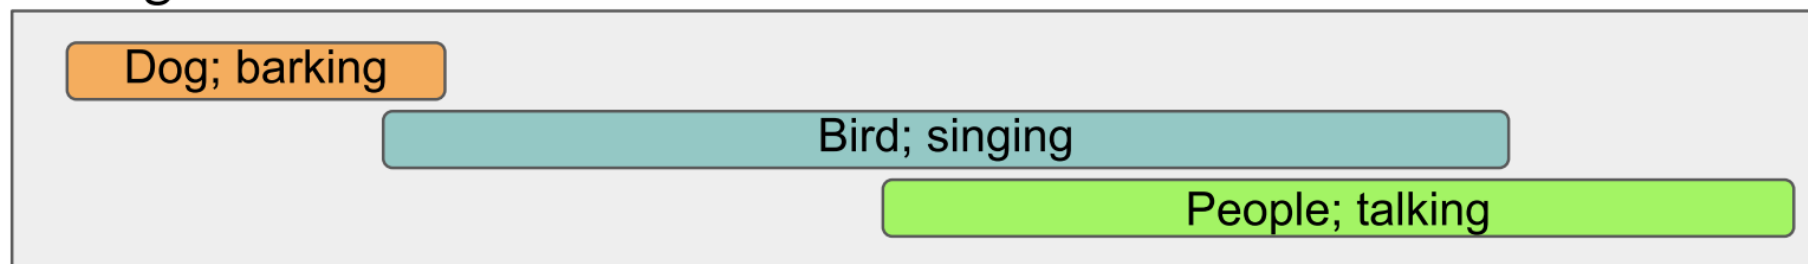
如有相关问题，建议请随时与我们联系！

- ▶ 语言模型
- ▶ 语音合成
- ▶ 图像生成
- ▶ **声音事件检测**
- ▶ 图片摘要生成
- ▶ 音视频场景识别

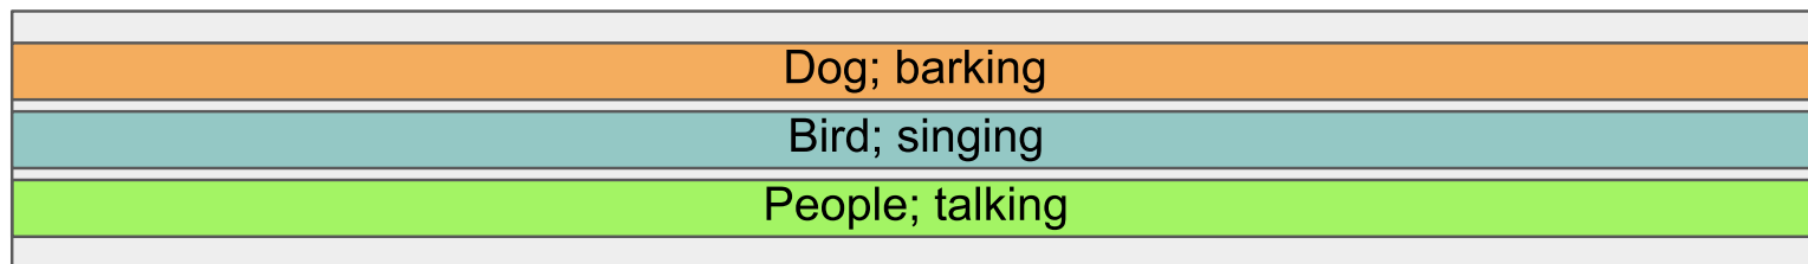
声音事件检测

- ▶ Weakly-labelled sound event detection
- ▶ Output: tagging (classification) and boundary detection

Strong labels

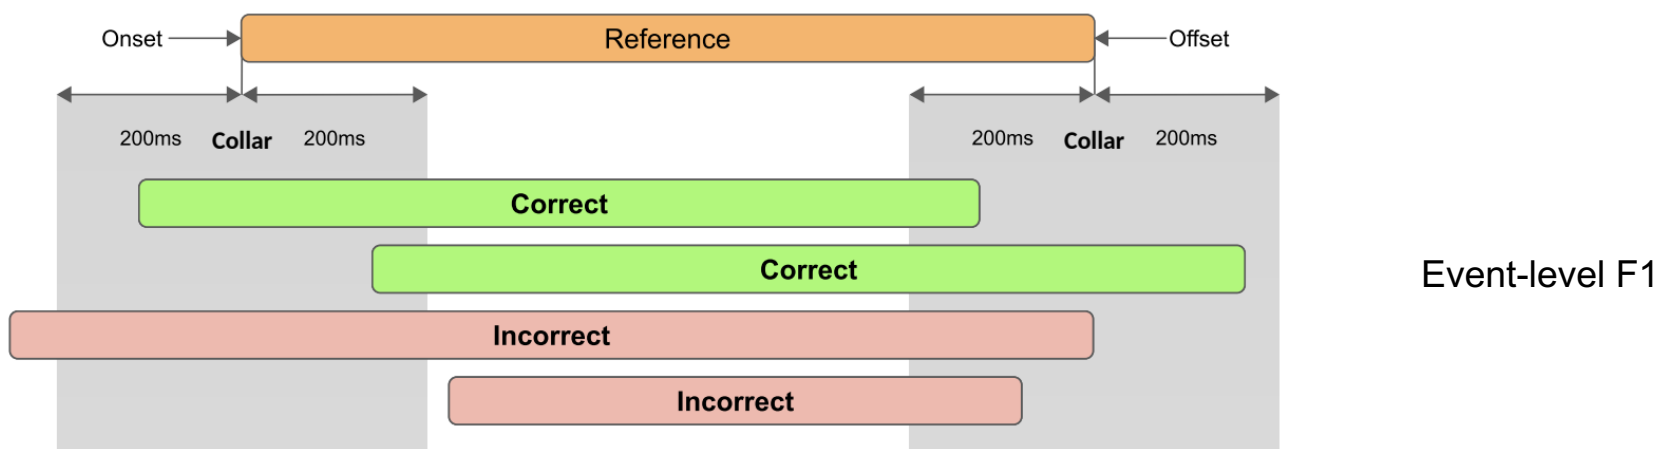



Weak labels



声音事件检测

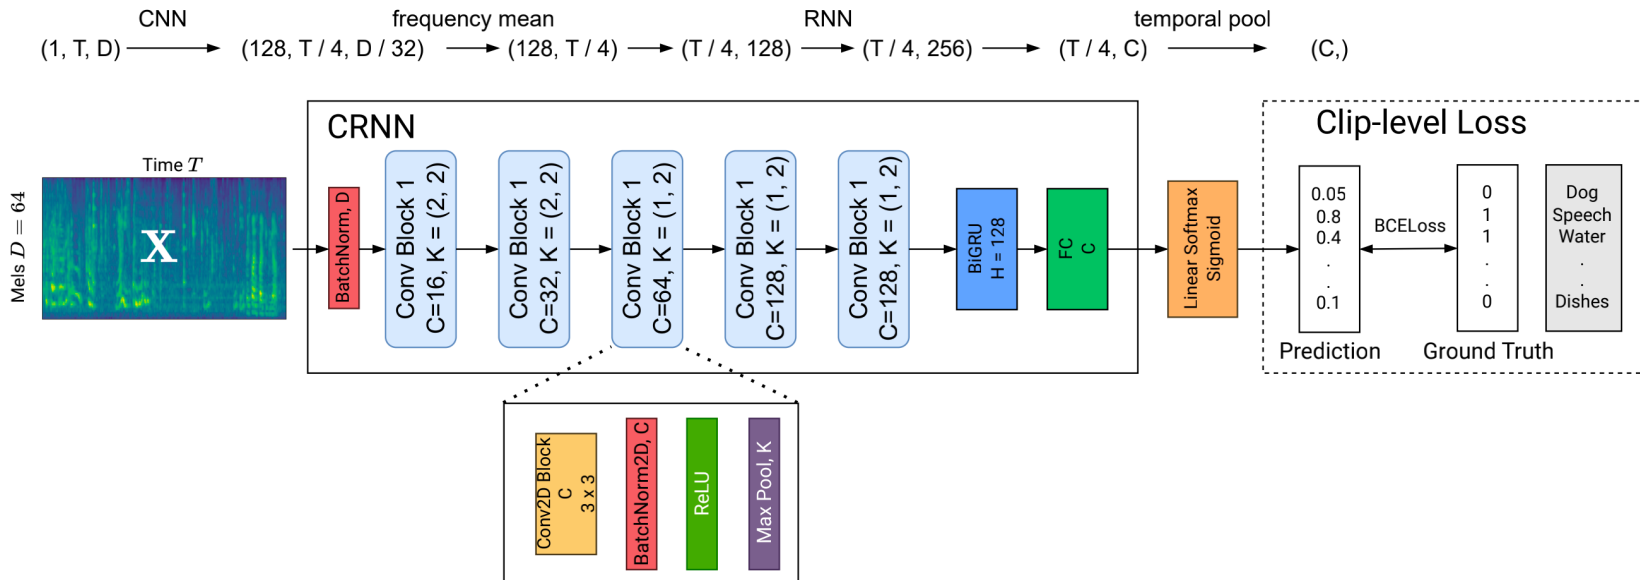
- Metrics: F1 score, Tagging evaluated by segment- and **event- F1** $F = \frac{2PR}{P + R}$



声音事件检测

- ▶ Data: DCASE18
- ▶ Baseline: CRNN
- ▶ Feature: LMS 64d
- ▶ Loss: Binary cross entropy (BCE):

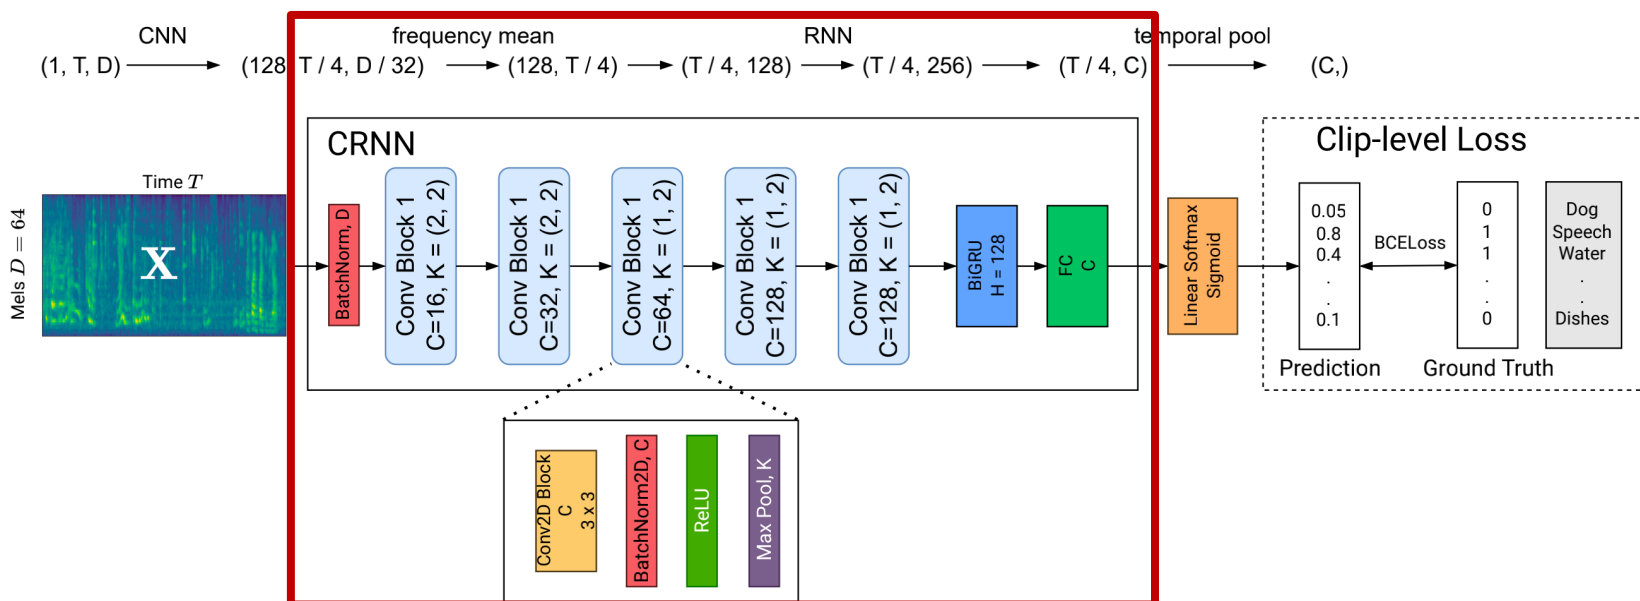
$$\mathcal{L}(y, \hat{y}) = -\hat{y} \log(y) + (1 - \hat{y}) \log(1 - y)$$



声音事件检测

- ▶ Data: DCASE18
- ▶ Baseline: CRNN
- ▶ Feature: LMS 64d
- ▶ Loss: Binary cross entropy (BCE):

$$\mathcal{L}(y, \hat{y}) = -\hat{y} \log(y) + (1 - \hat{y}) \log(1 - y)$$



► 基本要求

1. 理解代码逻辑，熟悉声音事件检测模型
2. 了解弱监督情况下进行时间轴预测的难点，以及基线模型设计的原理
3. 按模型框架实现CRNN模型

► 高阶要求

1. 可以修改模型深度、参数以及超参数，优化模型结果
2. 调研学习音频中的数据增强方法以应用

► 实践报告要求

1. 回答基本要求中的各个要点，若有额外工作，写明白修改的地方和额外的方法
2. 写出详细的实验过程和实验分析
3. 提交代码，给出重现最优结果的脚本和配置


▶ 参考文献

- ▶ Cao, Yin, Qiuqiang Kong, Turab Iqbal, Fengyan An, Wenwu Wang, and Mark D. Plumbley. "Polyphonic sound event detection and localization using a two-stage strategy." arXiv preprint arXiv:1905.00268 (2019).
- ▶ Mesaros, Annamaria, Toni Heittola, Tuomas Virtanen, and Mark D. Plumbley. "Sound event detection: A tutorial." IEEE Signal Processing Magazine 38, no. 5 (2021): 67-83.
- ▶ Dinkel, Heinrich, Mengyue Wu, and Kai Yu. "Towards duration robust weakly supervised sound event detection." IEEE/ACM Transactions on Audio, Speech, and Language Processing 29 (2021): 887-900.


- ▶ 语言模型
- ▶ 语音合成
- ▶ 图像生成
- ▶ 声音事件检测
- ▶ **图片摘要生成**
- ▶ 音视频场景识别

图片摘要生成

- ▶ A modality translation task
- ▶ Input: images; Output: natural language to describe the image
- ▶ Supervision signal: human-written captions, 5 captions/image



A little girl sitting on a bed with a teddy bear.




A group of people sitting on a boat in the water.



Evaluation: language generation metrics and image-language metrics

- ▶ BLEU4, 4-gram overlap precision
- ▶ METEOR: 衡量生成词语的 F1
 - ▶ 对词语做了 lemmatization (speaking -> speak) 并用 wordnet 映射到对应的概念
- ▶ CIDEr: 把句子表示成各个词语的 TF-IDF 组成的向量, 计算 reference 和 candidate 向量的 cosine similarity
- ▶ SPICE: 把句子转化成从 scene graph 中提取的 tuple set, 计算 reference 和 candidate tuple set 的 F1
- **SPIDEr:** (SPICE+CIDEr)/2

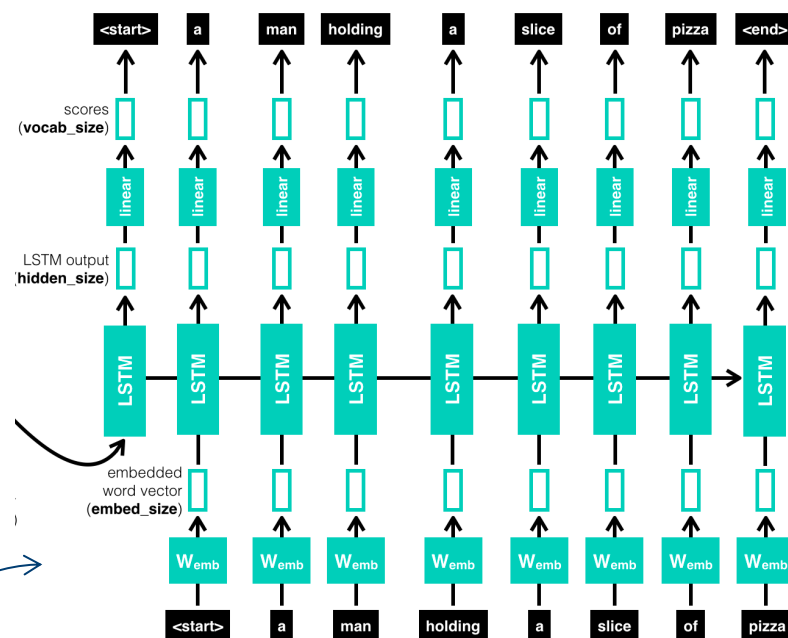
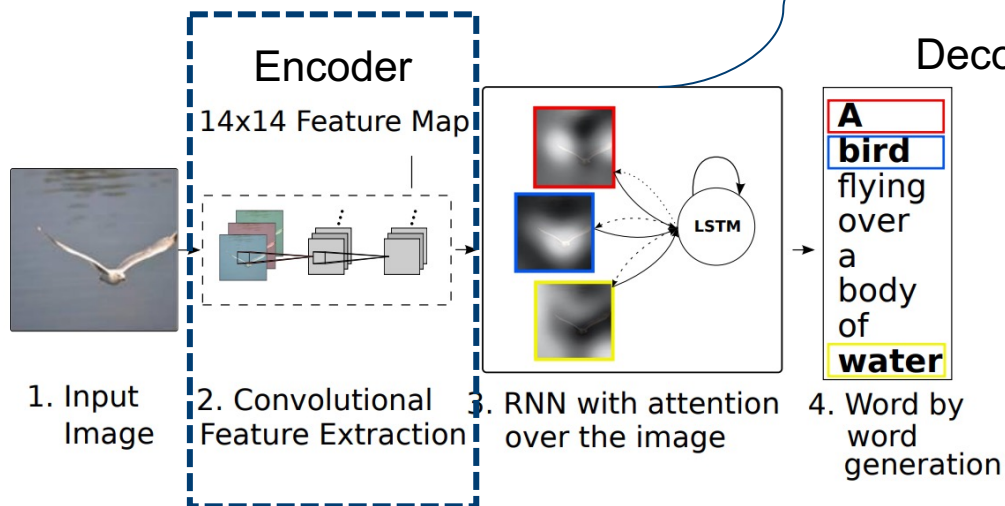
DT JJ amod NN nsubj VBG prep IN pobj NN prep IN DT NN nn NN
A young girl standing on top of a tennis court



{ (girl), (court), (girl, young), (girl, standing)
(court, tennis), (girl, on-top-of, court) }

图片摘要生成

- ▶ Dataset: Flickr8k
- ▶ Baseline: CNN-LSTM with attention
- ▶ Encoder output: $14 \times 14 \times 512$ feature map of the 4th convolutional layer before maxpool



► 基本要求

1. 理解跨模态模型对于不同模态数据处理的方式，编码器-解码器模型的整体框架
2. 实现scheduled sampling (知道在哪改，怎么改) *
3. 了解对于跨模态生成任务不同指标衡量的目的，打印生成最好和最坏的指标样例比较客观&主观评测

► 高阶要求

1. 通过修改代码和超参数进行模型调优，改进模型性能
2. 可以使用额外数据以及预训练模型

► 实践报告要求

1. 回答基本要求中的各个要点，若有额外工作增加详细做法说明
2. 写出详细的实验过程和实验分析，分析不同指标下的生成样例差异，与个人主观评测进行对比
3. 提交代码，给出重现最优结果的脚本和配置

*Scheduled sampling: Bengio, Samy, Oriol Vinyals, Navdeep Jaitly, and Noam Shazeer. "Scheduled sampling for sequence prediction with recurrent neural networks." Advances in neural information processing systems 28 (2015).

▶ 参考文献

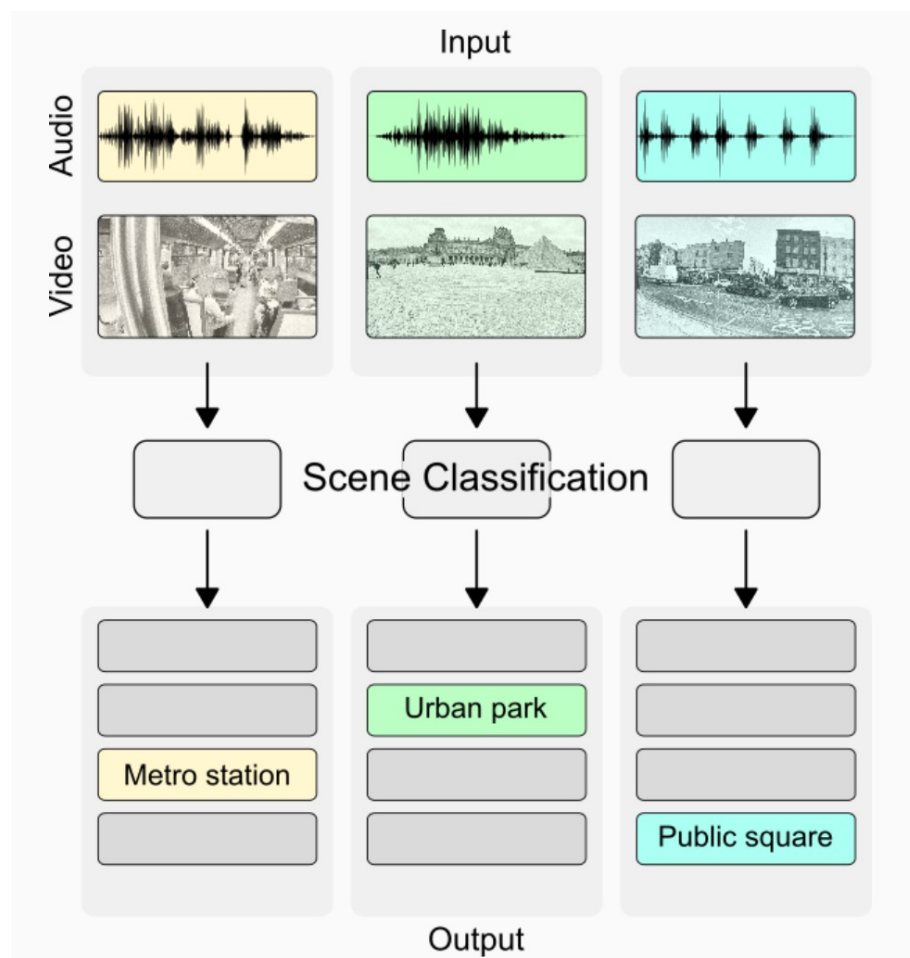
- ▶ Xu, Kelvin, Jimmy Ba, Ryan Kiros, Kyunghyun Cho, Aaron Courville, Ruslan Salakhudinov, Rich Zemel, and Yoshua Bengio. "Show, attend and tell: Neural image caption generation with visual attention." In International conference on machine learning, pp. 2048-2057. PMLR, 2015.
- ▶ Wang, Haoran, Yue Zhang, and Xiaosheng Yu. "An overview of image caption generation methods." Computational intelligence and neuroscience 2020.
- ▶ You, Quanzeng, Hailin Jin, Zhaowen Wang, Chen Fang, and Jiebo Luo. "Image captioning with semantic attention." In Proceedings of the IEEE conference on computer vision and pattern recognition, pp. 4651-4659. 2016.

- ▶ 语言模型
- ▶ 语音合成
- ▶ 图像生成
- ▶ 声音事件检测
- ▶ 图片摘要生成
- ▶ **音视频场景识别**

音视频场景识别

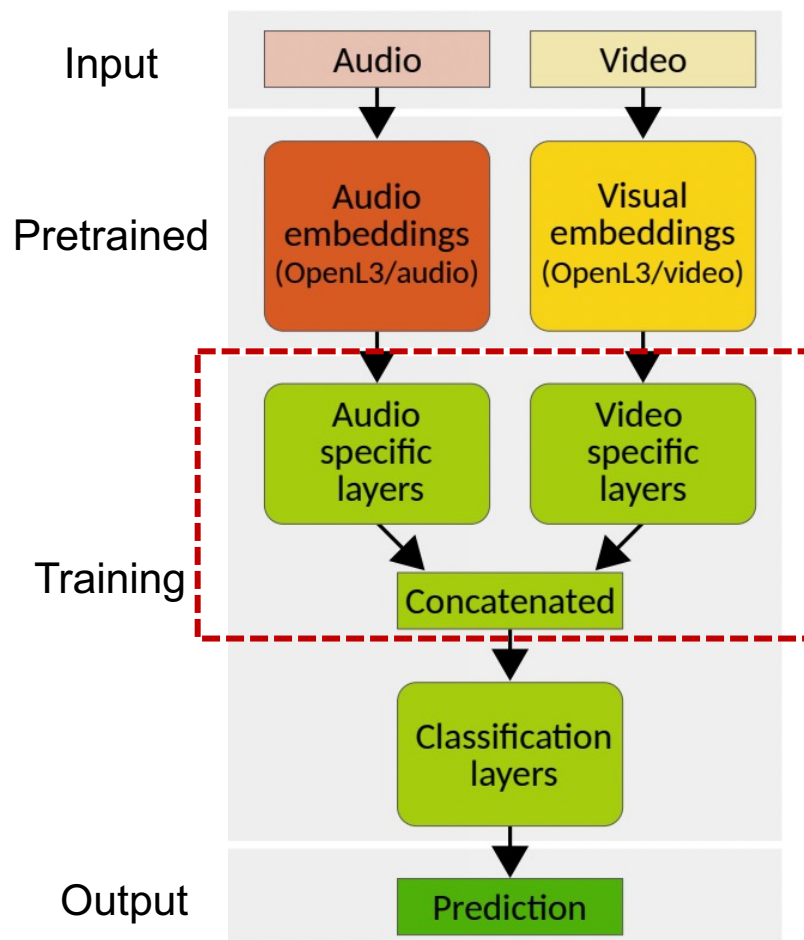

- ▶ Multi-modal representation learning
- ▶ Input: video and audio
- ▶ Output: a uni-label scene classification result

1. Airport
2. Indoor shopping mall
3. Metro station
4. Pedestrian street
5. Public square
6. Street with medium level of traffic
7. Travelling by a tram
8. Travelling by a bus
9. Travelling by an underground metro
10. Urban park



音视频场景识别


- ▶ Pretraining can be a useful tool for multi-modal representation: baseline utilized pretrained audio and video embeddings
- ▶ Feature concat is the most basic modality fusion method, however we can also involve separate audio and video networks first to learn task-specific representations



音视频场景识别

- ▶ Evaluation metrics:
 - ▶ Audio-only, Video-only, A/V combined
 - ▶ Class-wise Results
 - ▶ LogLoss

$$L_{\log}(y, p) = -(y \log(p) + (1 - y) \log(1 - p))$$



► 基本要求

1. 理解多模态融合工作不同表征获取的方式、模态融合的方式
2. 分析有冲突的模态结果, i.e. 在何种类别上多模态融合更有效, 为什么
3. 替换为early特征融合和late决策融合的两类融合方式

► 高阶要求

1. 通过替换特征、修改模型、和超参数进行模型调优, 改进模型性能

► 实践报告要求

1. 回答基本要求中的各个要点, 若有额外工作提供详细说明
2. 写出详细的实验过程和实验分析
3. 提交代码, 给出重现最优结果的脚本和配置

▶ 参考文献

- ▶ <https://dcase.community/challenge2021/task-acoustic-scene-classification#subtask-b>
- ▶ Wang, Shanshan, Toni Heittola, Annamaria Mesaros, and Tuomas Virtanen. "Audio-visual scene classification: analysis of DCASE 2021 Challenge submissions." arXiv preprint arXiv:2105.13675 (2021).
- ▶ Wang, Shanshan, Annamaria Mesaros, Toni Heittola, and Tuomas Virtanen. "A curated dataset of urban scenes for audio-visual scene analysis." In ICASSP 2021-2021 IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP), pp. 626-630. IEEE, 2021.

课程作业问题联系人

授课教师： 吴梦玥 邮箱： mengyuewu@sjtu.edu.cn
▶ 助教： 谢泽宇 邮箱： zeyu_xie@sjtu.edu.cn

Github QA page:

https://github.com/chensexie95/deeplearning_course_sjtu/issues

按疑问需求设置线上/线下统一答疑时间

- ▶ 语言模型、语音识别、VAE图片生成
 - ▶ 代码: https://github.com/chenxie95/deeplearning_course_sjtu
 - ▶ 数据: /lustre/home/acct-stu/stu281/deeplearning_course_sjtu
- ▶ 音视频场景识别、声音事件检测、图片摘要生成:
 - ▶ 代码: /lustre/home/acct-stu/stu282/Project/{av_scene_classify, sound_event_detection, image_captioning}
 - ▶ 数据: /lustre/home/acct-stu/stu282/Data/{av_scene_classify, sound_event_detection, image_captioning}
 - ▶ 环境: “conda env create -f /lustre/home/acct-stu/stu282/Project/env.yaml”

“交我算” 平台使用

▶ 登录节点

- ▶ 使用 SSH 登录:

```
[local] $ ssh [username]@pilogin.hpc.sjtu.edu.cn
```

- ▶ 写入 ~/.ssh/config

```
Host pi
```

```
    HostName pilogin.hpc.sjtu.edu.cn
```

```
    User [username]
```

```
    ServerAliveInterval 240
```

登录命令:

```
[local] $ ssh pi
```

- ▶ 免密登录 (Linux / Mac)

- ▶ 生成私钥: [local] \$ ssh-keygen

- ▶ 发送私钥到登录节点: [local] \$ ssh-copy-id pi

“交我算” 平台使用

▶ 数据传输

- ▶ 基本命令: `$ scp [source] [destination]`

- ▶ 例: 从超算上拷贝数据到本地:
`[local] scp pi:~/data.txt ./`

▶ 终端复用: Tmux

- ▶ 建立新的 session: `[pi] tmux new -s [session_name]`
- ▶ 暂离session: `Ctrl+b d` (先同时按Ctrl和b, 再按d, 下同)
- ▶ 重新进入session: `[pi] tmux a -t [session_name]`
- ▶ 切换session: `Ctrl+b s`

“交我算” 平台使用

▶ 用 Slurm 系统提交作业

▶ 交互式作业 `srun` / `salloc`:

- ▶ `srun`: `[pi] $ srun -N 1 -p small --pty /bin/bash`
- ▶ `salloc`: `[pi] $ salloc -N 1 -p small`
- ▶ 随后登录分配到的节点

▶ 提交作业脚本 `sbatch`:

▶ 写好脚本, 如 `run.sh`:

```
#!/bin/bash
#SBATCH --job-name hostname
#SBATCH -p small
#SBATCH -N 1
/bin/hostname
```

- ▶ 提交任务: `[pi] $ sbatch run.sh`
- ▶ 该作业与 `[pi] $ srun --job-name hostname -p small -N 1 /bin/hostname` 一致

▶ 提交 gpu 任务:

- ▶ 指定队列: `-p dgx2`
- ▶ 申请一块GPU: `--gres gpu:1`

“交我算” 平台使用

使用注意事项:

- ▶ 为避免计算资源浪费, 教学账号限制作业运行数量1个、核数10个、GPU卡数1卡、最长运行时间12小时。请不要在登录节点运行作业, 否则将会被封禁!!! 教学支撑gpu队列为dgx2 (单卡拥有32G显存)。目前集群GPU资源紧张, 可能会出现排队现象, 请妥善安排作业提交时间, 并且不要运行与课程无关的任务。
- ▶ 集群状态查询: <https://status.hpc.sjtu.edu.cn/>

相关文档:

- ▶ 登录: <https://docs.hpc.sjtu.edu.cn/login/index.html>
- ▶ 作业提交: <https://docs.hpc.sjtu.edu.cn/job/index.html>

如有任何问题, 请联系助教。

课程论坛: https://github.com/chensexie95/deeplearning_course_sjtu/issues